

Jiewen_Luo_ps1

Jiewen

1/17/2020

1. Statistical and Machine Learning

Supervised machine learning is an algorithm where a mathematical model is built from a set of labeled data contain both inputs X s and the desired output Y . In supervised machine learning, the system needed to be trained to perform certain tasks. Unsupervised machine learning is a self-learning process. Under this system, an algorithm is given only the unlabeled input data X s without corresponding output variables.

By building an appropriate model that establishes the corresponding relationship between X s and Y , supervised learning can predict output data Y when new X s in the data set is introduced. To accomplish this task, the algorithm iteratively predicts the training data until a given level of predicting accuracy is achieved. As for unsupervised learning, the purpose is to summarize and explain the underlying distribution or structure of the data.

The actual data generating processes in supervise learning also differ from unsupervised learning. To solve any given supervised learning problem, we first need to determine the types of data needed. The next step is to collect representative training data that well-represented the entire population. Afterward, the input feature representation of the learned function should be well-selected. A good feature vector should contain enough information for predicting output but not too large because of the course of dimensionality. Another critical step is to select the structure of the learning function and learning algorithm. Finally, the optimal model can establish after iteratively running the learning algorithm on the collected training data. For unsupervised learning, the training step is eliminated so the system has to dig in hidden patterns of the data on its own.

In terms of the methods applied, these two types of machine learning systems also distinct from each other. When dealing with continuous quantitative data, supervised learning performs regression analysis while unsupervised learning use dimensionality. When discrete data or qualitative data are given, supervised learning handles the task by classification whereas unsupervised learning adopts cluster analysis.

Here are two machine learning tasks that help to illustrate the difference between these two algorithm systems. When dealing with the grouping issue, supervised learning applies a classification algorithm. Because the data given come with categories, the task for the system is to accurately recognize category membership of each observation in the training set. After the training, the system is then able to identify the categories of new observations. In the unsupervised case, the categories of the data are unknown. Therefore, cluster analysis is used to segment data into groups. This goal can be achieved by

identifying the degree of inherent similarity in the data based on the presence or absence of various indicators. In the task of density estimation, the difference can also be presented mathematically. While supervised learning aims at inferring a conditional probability distribution conditioned on the label x of input data, unsupervised learning aims at inferring a priori probability distribution.

Finally, from an evaluation perspective, these two algorithms vary in terms of complexity and accuracy. Supervised learning is a simpler method compared to unsupervised learning. In supervised learning, because the model built by supervised learning is repeatedly examined by the training data set until it reaches a target level of correctness within the training data, the result is trustworthy. However, there is no such equivalent model optimization process in unsupervised machine learning, therefore, the result is more doubtful.

2. Linear Regression

a.

```
mc=mtcars
Predict1=lm(mpg~cyl,data=mc)
summary(Predict1)

##
## Call:
## lm(formula = mpg ~ cyl, data = mc)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.9814 -2.1185  0.2217  1.0717  7.5186
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  37.8846     2.0738   18.27 < 2e-16 ***
## cyl         -2.8758     0.3224   -8.92 6.11e-10 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.206 on 30 degrees of freedom
## Multiple R-squared:  0.7262, Adjusted R-squared:  0.7171
## F-statistic: 79.56 on 1 and 30 DF, p-value: 6.113e-10
```

Output value: F test shows that the model is statistically significant. R squared is high, suggesting that the model is quite a good fit. Parameter value for **cylinders** is statistically significant with value of -2.8758.

b. $\widehat{mpg}=37.8846-2.8758*cyl$;

c.

```
Predict2=lm(mpg~cyl+wt,data=mc)
summary(Predict2)

##
## Call:
## lm(formula = mpg ~ cyl + wt, data = mc)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.2893 -1.5512 -0.4684  1.5743  6.1004
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  39.6863     1.7150   23.141 < 2e-16 ***
## cyl         -1.5078     0.4147   -3.636 0.001064 **
## wt          -3.1910     0.7569   -4.216 0.000222 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.568 on 29 degrees of freedom
## Multiple R-squared:  0.8302, Adjusted R-squared:  0.8185
## F-statistic: 70.91 on 2 and 29 DF,  p-value: 6.809e-12
```

Now $\widehat{mpg}_2 = 39.6863 - 1.5078 \cdot cyl - 3.1910 \cdot wt$.

By adding **vehicle weight** into the model, the coefficient size of the constant term increases but the coefficient size for variable **cylinders** decreases. This means that the negative effect of **cylinders** on **miles per gallon** decreases. Additionally, this model also shows us that **vehicle weight** has statically and substantively significant negative effect on **miles per gallon**.

d.

```
mc=mtcars
Predict3=lm(mpg~cyl+wt+cyl:wt,data=mc)
summary(Predict3)

##
## Call:
## lm(formula = mpg ~ cyl + wt + cyl:wt, data = mc)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.2288 -1.3495 -0.5042  1.4647  5.2344
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  54.3068     6.1275   8.863 1.29e-09 ***
## cyl         -3.8032     1.0050   -3.784 0.000747 ***
## wt          -8.6556     2.3201   -3.731 0.000861 ***
```

```
## cyl:wt      0.8084      0.3273      2.470 0.019882 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2.368 on 28 degrees of freedom
## Multiple R-squared:  0.8606, Adjusted R-squared:  0.8457
## F-statistic: 57.62 on 3 and 28 DF,  p-value: 4.231e-12
```

$\widehat{mpg}_2 = 54.3068 - 3.8032 \cdot cyl - 8.6556 \cdot wt + 0.8084 \cdot wt \cdot cyl$. After adding interaction term, the joint effect of these three variable terms is statistically significant. We can also see that the intercept term remains positive, and the estimates for **cylinders** and **vehicle wage** remain negative. However, the magnitude of the estimator for **cylinders** and **vehicle wage** both increase.

By adding interaction term, we are theoretically asserting that **vehicle wage** influences the relationship between **cylinders** and **miles per gallon**.

3. Non-linear Regression

a.

```
dt=read.csv("wage_data.csv")
model1<-lm(wage~poly(dt$age,2),data=dt)
summary(model1)

##
## Call:
## lm(formula = wage ~ poly(dt$age, 2), data = dt)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -99.126 -24.309  -5.017   15.494  205.621
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    111.7036     0.7302   152.99  <2e-16 ***
## poly(dt$age, 2)1  447.0679     39.9926    11.18  <2e-16 ***
## poly(dt$age, 2)2 -478.3158     39.9926   -11.96  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 39.99 on 2997 degrees of freedom
## Multiple R-squared:  0.08209,    Adjusted R-squared:  0.08147
## F-statistic:   134 on 2 and 2997 DF,  p-value: < 2.2e-16
```

$\widehat{wage}_1 = 111.7036 + 447.0679 \cdot age - 478.3158 \cdot age^2$


The model is statistically significant, but has poor goodness of fit. The linear age term has statistically and substantively significant positive value and the quadratic age term has statistically and substantively significant negative value. This means that age has a diminishing marginal effect on wage. The marginal effect is positive at the beginning, but as age increases, the marginal effect will diminish to 0 and become negative afterward.

b.

```
library(MASS)

## Warning: package 'MASS' was built under R version 3.6.2

plot(sort(dt$age), fitted(model1)[order(dt$age)], xlab="Age", ylab="Wage",
     col="green", type='l')
preds <- predict(model1, data.frame(x=dt$age), interval =
"confidence", level=0.95)
x<-sort(dt$age)
y1<-preds[,3][order(dt$age)]
y2<-preds[,2][order(dt$age)]
polygon(c(x, rev(x)), c(y2, rev(y1)), col = 'grey', border = NA)
lines(sort(dt$age), preds[,1][order(dt$age)], lty = 'solid', col = 'green')
lines(sort(dt$age), preds[,3][order(dt$age)], lty = 'dashed', col = 'red')
lines(sort(dt$age), preds[,2][order(dt$age)], lty = 'dashed', col = 'blue')
legend("bottom", c("fitted", "upper bound", "lower bound"),
      col=c("green", "red", "blue"), lwd=3)
```



c.

From the graph we can see that the 95% confident interval is narrowing down from age 20 to approximately age 35 but gradually expanding afterward.


By fitting a polynomial regression, we are asserting a non-linear relationship between our predictor and predictee.

d.

```
model2<-lm(dt$wage~dt$age,data=dt)
summary(model2)

##
## Call:
## lm(formula = dt$wage ~ dt$age, data = dt)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -100.265  -25.115   -6.063   16.601   205.748
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  81.70474     2.84624   28.71  <2e-16 ***
## dt$age        0.70728     0.06475   10.92  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 40.93 on 2998 degrees of freedom
## Multiple R-squared:  0.03827,    Adjusted R-squared:  0.03795
## F-statistic: 119.3 on 1 and 2998 DF,  p-value: < 2.2e-16

plot(sort(dt$age), fitted(model1)[order(dt$age)],xlab="Age",ylab="Wage",
col="green", type='l')
lines(sort(dt$age), fitted(model2)[order(dt$age)],col="yellow")
```



The linear regression function is: $\widehat{wage2} = 81.70474 + 0.70728 \cdot age$.

This linear regression model is statistically significant, but the goodness of fit is even worse than the polynomial model. Age linear term is still statistically significant at 99% confident interval, but now age as constant positive effect on wage. More specifically, one unit increase in age will lead to increase in wage by 0.70728 unit.